

Gods True Gospel

A CONCORDANT BIBLE-STUDY SKILL



SOLA SCRIPTURA

HEBREW · ARAMAIC · GREEK

Preface

This document describes the Gods True Gospel skill: what it is, how it works, what it is built from, and why it has been engineered the way it has. The skill is a research toolkit for studying the Bible directly from its Hebrew, Aramaic and Greek source texts. It uses a concordant translation discipline and a typology corpus that ties patterns to verses rather than to schools of interpretation.

The intended reader is a developer, researcher, or technically-minded student of Scripture who wants to understand both the methodology and the mechanics. Sections alternate between functional explanation, technical reference, and architectural rationale. Diagrams accompany each major idea so the document can be read either linearly or by jumping to the figure that matches your question.

Throughout this documentation, one principle is held without exception: every claim the skill produces must be traceable to the source text itself. No church father, doctrinal school, or training-data echo enters the answer. That principle, sola scriptura, is not a slogan in this project. It is a rule that shapes every script, every entry, and every output block. Its consequences for design, for translation methodology, and for the pattern-discovery process are described in the chapters that follow.

Conventions. Hebrew, Aramaic and Greek words appear in transliteration where possible to support readers without those scripts installed. Strong codes use the standard *Hxxxx* for Hebrew or Aramaic and *Gxxxx* for Greek. Verse references follow the form *Book Chapter:Verse*, for example Rom 5:14.

Contents

Preface	2
1. Introduction	3
2. Architectural principles	5
3. The concordant translation methodology	8
4. Functional walkthrough	11
5. Technical implementation	14
6. The knowledge corpus	16
7. The typology framework	18
8. Repository layout	20
9. Quality assurance	21
Appendix A. Glossary	22
Appendix B. References and attribution	24
Appendix C. Sample workflows	25

1. Introduction

What problem does this project solve, and what does it offer that other Bible-study tools do not?

1.1 The problem

Most Bible-study software answers questions about Scripture by pulling from one or more popular translations and combining that with an interpretive layer drawn from commentary traditions. The answer is shaped twice before it reaches the reader: first by the choices of the translation team, and second by the doctrinal lens of the commentary tradition. When a study question concerns the linguistic continuity of a single source word (where the same Hebrew or Greek word appears across the canon, in what contexts, and how its meaning is built up) the layered reshape obscures the very thing the question is about.

Modern AI assistants add a third layer: the model's training data. Even when a model is asked to be neutral, it will, by default, fall back on the most-represented interpretive traditions in its training corpus.

1.2 The skill in one paragraph

Gods True Gospel is a research toolkit that produces text-grounded answers about Scripture by reading the Hebrew, Aramaic and Greek source text directly from a local corpus and applying a concordant translation discipline plus a 170-entry typology framework. Every claim is traceable to a verse and to a Strong-coded source word. No church father, midrash tradition, doctrinal school, or modern interpreter is consulted. No training-data echo is used.

1.3 Sola scriptura, in this project

The phrase *sola scriptura* has a long history and many meanings. In this project it has a specific, technical meaning: the operational rule that no claim in the output may rest on anything other than the source text bundled in Kennis/puur/ and Kennis/strong/.

Prohibited as primary input: church fathers, midrashic commentary, rabbinic tradition, modern typology authors, doctrinal schools, training-data echoes, translation tradition (KJV, NBG-1951, Reina-Valera, and others) used as authority rather than as comparison reference.

Required for any claim: a verse reference plus a citation of the source-text form (Hebrew, Aramaic, or Greek) with transliteration and Strong code. For a typological claim, two independent witnesses in the canon are required. For a contested claim, the skill must mark it as N3 (suggested, not proven) rather than overstate.

1.4 What you can do with the skill

Three kinds of inquiry are supported. **Word-study:** given a Hebrew or Greek word, find every occurrence in the canon and read a concordant rendering. **Verse-study:** given a verse, retrieve the source text with per-word Strong tagging and identify activated typology entries. **Pattern-study:** given a hypothesis, inventory all occurrences in the source text, group by meaning-cluster, and confirm via the two-witness rule.

2. Architectural principles

Six principles shape the architecture. Each is engineered into the codebase so that the principle is checkable, not merely declared.




Figure 1 — Top-level component diagram. The skill never serves answers from training memory; every block reaches into the corpus.

2.1 Source-text discipline

All linguistic claims in the system are anchored in the local source-text corpus. Kennis/puur/ contains the source text per Bible book in line-delimited JSON with transliteration. Kennis/strong/ contains the same text with per-word Strong-codes plus parsing information. Both layers cover all 70 books.

2.2 No external interpretive authority

The corpus is intentionally bare of interpretive tradition. There is no patristic commentary file, no midrash extract, no doctrinal-school summary. This is a design choice.

2.3 The two-witness rule for typology

A pattern is admitted to the typology corpus only when at least two independent witnesses in the canon support it, and when the pattern can be verified by reproducible search scripts. The rule comes from Deuteronomy itself: a matter is established by the mouth of two or three witnesses (Deu 19:15).

2.4 Reproducibility by script

Every claim in a typology entry must be reproducible by a search script. Each entry contains a `zoekquery_gebruikt` field that records the exact command used. Running the same command

tomorrow on the same corpus must yield the same verse list.

2.5 The attentive-reader presumption

The Bible was written for readers who knew the text well. The skill therefore reads for design rather than for accident. When a structural feature appears that looks deliberately placed (a number repeated across narratives, a phrase echoed across testaments, a chiasmic shape spanning chapters) the skill treats it as design unless the textual evidence specifically rules it out.

2.6 The direct reading wins on conflict

Typology and structural patterns serve as a coherence check, never as an override. When a typological pattern conflicts with a direct reading, the conflict is reported to the user as a signal, but the direct reading wins.

2.7 Engineered consequences

These principles have concrete engineering consequences. The corpus is local and offline. The preflight uses only regex against a static trigger table. The typology entries are flat Markdown files. The skill is fully self-contained.

3. The concordant translation methodology

Why this project translates the way it does, and how the chosen renderings differ from the choices made by major Bible translations.

3.1 What is concordant translation?

Concordant translation is a discipline that renders one source-language word with a single fixed target-language word across all of its occurrences. If the Greek word *aiōn* appears in 122 places, all 122 places use the same target word. The trade-off is that smooth target-language prose is sacrificed in exchange for transparency about what the source actually keeps repeating.

3.2 Why we chose concordance

If the skill must trace every claim to the source text, the user must be able to follow what the source text actually says. If a single source word is rendered with three or four different target words, the linguistic continuity becomes invisible. Repeated motifs disappear. Word-plays evaporate. Cross-passage echoes break.


There is a second motivation. Every Bible translation carries doctrinal commitments that influence word choice. Concordance forces the same word everywhere, which forecloses one of the easier paths to doctrinal coloring.

3.3 How the discipline operates

When the skill produces a Block D response, it consults the concordant master for the user's target language. Where the chosen equivalent differs sharply from the rendering in a major reference translation (KJV in English, NBG-1951 in Dutch, Reina-Valera 1960 in Spanish), the divergence is named explicitly.

Concordant translation discipline

One source word renders consistently. The reader sees the linguistic continuity.



3.4 The skill thinks in source language

Internally, the skill reasons about the source-language form of a word, not about its translation. When a question mentions "forever" in English, the skill identifies the underlying Greek *aiōn*, *aiōnios*, or the Hebrew *olam*. The user therefore sees, side by side, what the source text says with one word and what the familiar translation rendered it as.

3.5 Worked examples of divergence

aiōn (G165, 122 occurrences). Concordant: "eon". Tradition renders "forever", "eternity", "world", or "age" depending on context. The concordant rendering preserves temporal scope. Tradition often implies infinite duration where the source marks finite scope.

hamartia (G266, 173 occurrences). Concordant: "miss-the-mark" preserves the archery metaphor. Tradition renders "sin", which conflates several distinct Greek terms (*hamartia*, *parabasis*, *paraptōma*, *anomia*).

nephesh (H5315, 754 occurrences). Concordant: "soul", applied also to animals. Tradition switches to "life", "creature", or "person" for animals.

ekklēsia (G1577, 114 occurrences). Concordant: "called-out" or "assembly". Tradition renders "church", which misframes Acts 19:32 (a riotous mob in a theater).

3.6 Three target languages, same discipline

The masters cover Dutch, English, and Spanish. The Spanish master fixes *aiōn* as "eón" rather than Reina-Valera "siglo" or "eternidad". The Dutch master fixes "aion" rather than NBG-1951 "eeuw". The English master fixes "eon" rather than KJV "forever".

3.7 What the discipline does not do

Concordance does not eliminate interpretation. The choice of which target word to fix is itself an interpretive act. The discipline forces that act into the open: once per source word, with a recorded rationale.

4. Functional walkthrough

What happens when a user asks the skill a question.

4.1 Three operating modes

Passive mode is the default. **Active mode** is used during research sessions for new typology entries. **Cross-check mode** is the audit mode that scans the corpus for inconsistencies.

4.2 From question to answer




Figure 3 — A user question becomes a response only after passing through eight stages, each rooted in the corpus.

4.3 Stage detail

Stage 1. Parse question, detect language, extract verse references. **Stage 2.** Preflight detection against three trigger tables. **Stage 3.** Source lookup from Kennis/strong/ with full Strong-tagging. **Stage 4.** Typology load for activated entries. **Stage 5.** Concordant pass, applying the master for the user's language. **Stage 6.** Synthesis: build Block A (direct reading), Block B (source citations), Block C (typology cross-references). **Stage 7.** Coherence check with N1/N2/N3 marking. **Stage 8.** Emit response in the user's language.

4.4 Block-structured response

The response is built from named blocks. Each block has a fixed purpose. Users can request the full dossier (Block A through E), a subset, or a single block.

Block-structured response (Block A-F)

Every skill response is built from named blocks; the user can ask for any subset.

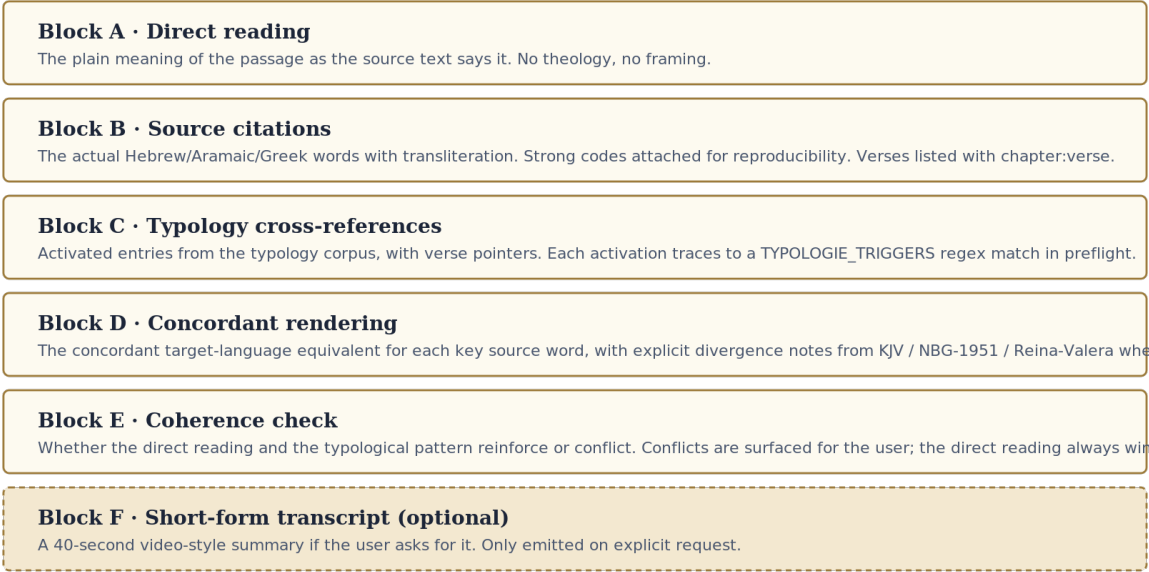


Figure 4 — Each block can be requested or suppressed individually. The full dossier (A-E) is the default.

4.5 Sample interaction

A user asks: *What does the source text say about the first and last Adam in 1 Corinthians 15?*

The preflight matches both A_entiteit/adam and H_contrast/eerste-laatste-adam. Stage 3 loads 1 Cor 15:21-22 and 15:45-49. Stage 5 surfaces the concordant rendering for key Greek words: *prōtos*, *eschatos*, *choikos*, *epouranios*. Stage 7 marks the entire response as N1 because Romans 5:14 explicitly calls Adam *typos tou mellontos*.

5. Technical implementation

The scripts that drive the skill, the data formats they consume, and the deterministic logic.

5.1 The skill core scripts

Script	Purpose
skill_v5_preflight.py	Question parsing, trigger detection, mode selection.
skill_v5_lookup.py	Verse and Strong-code lookup.
skill_v5_blocks.py	Block A through E construction.
skill_v5_transcript.py	Block F transcript generation.
skill_v5_interview.py	Active-mode protocol.
typologie_zoek.py	Source-text search helper.

5.2 Preflight detection

The preflight matches the question against per-entry triggers, category triggers, and external-tradition signals. Example:

```
{
  'H_contrast/eerste-laatste-adam': {
    'patterns': [r'\beerste.{0,10}(adam|laatste)', r'\blaatste\s+adam\b'],
    'pad': 'Kennis/typologie/H_contrast/eerste-laatste-adam.md',
  }
}
```

5.3 The source-text search helper

typologie_zoek.py accepts several search modes: by Strong code, number-pattern, word-root, day-pattern. Output is JSON-readable and deterministic.

```
python3 skill/typologie_zoek.py --strong G5179
python3 skill/typologie_zoek.py --cijfer 120
python3 skill/typologie_zoek.py --dag 3
```

5.4 Reverse lookup index

Kennis/typologie/_index.json maps every verse reference and every Strong code to the list of typology entries that cite it. The index covers 170 entries, approximately 3300 unique verse references, and approximately 3300 unique Strong codes.

5.5 Cross-reference checker and data formats

check_xrefs.py audits the corpus for internal-link integrity. Three formats: JSONL for source text, JSON for indexes and masters, Markdown for typology entries.

5.6 No proprietary stack

The skill runs on Python 3.10 or higher. No proprietary database, no commercial Bible software, no external API.

6. The knowledge corpus

The corpus is organized in three layers, ordered from raw to abstract.


Layer	Path	Content	Format
1. Raw	Kennis/puur/	Source text per Bible book, transliterated.	JSONL
2. Tagged	Kennis/strong/	Source text + per-word Strong codes + parsing	JSONL
3. Depth	Kennis/diepte/	Word studies per Strong code.	Markdown
Indexes	Kennis/index/	Reverse Strong to verses.	JSON
Mappings	Kennis/lxx-mapping-*	Hebrew to Greek (LXX).	JSON
Masters	Kennis/masters/	Concordant masters (NL, EN, ES).	JSON+MD
Protocols	Kennis/protocollen/	Methodological protocols.	Markdown
Typology	Kennis/typologie/	170 entries.	Markdown

The raw text covers 70 books (Hebrew Bible as 39 books, Greek New Testament as 27, Psalms as five sub-files). The Strong-tagged text adds two fields per word. The depth studies follow a seven-anchor template. The masters cover three languages with rationale per word.

7. The typology framework

The three-axes typology framework

Each entry is positioned simultaneously on all three axes



Typology has historically been a soft discipline. This project treats it as hard. Each entry is positioned along three axes: **category** (eight major categories A through H), **hermeneutic layer** (L1 direct-historical, L2 Christ-typological, L3 ecclesiological, L4 eschatological), and **certainty** (N1 hard, N2 inferable, N3 suggested). The corpus has 170 entries: 74 entity, 13 number, 16 time, 8 language, 15 structure, 17 narrative, 15 role, 12 contrast.

Each entry follows the V2 template with twelve required sections and passes the seven-step protocol.

8. Repository layout

```
github/  
├── README.md  
├── LICENSE  
├── ATTRIBUTION.md  
├── CONTRIBUTING.md  
├── Kennis/ (puur, strong, diepte, index, masters, protocollen, typologie)  
├── skill/ (skill_v5_*.py, fase10/, typologie_zoek.py)  
├── docs/ (godstruegospel-documentation.pdf)  
├── build-software/  
└── Testen/
```


9. Quality assurance

How correctness is established and maintained.

9.1 Pairwise testset

Quality is established through a pairwise testset documented in Testen/skill-v5-pairwise-testset.md. The testset has 25 cases across four validation dimensions: output-type correctness, source-discipline, edge cases, and end-to-end. Round reports record performance across iterations.

9.2 Test gate per phase

A development phase is not closed until the phase's testset passes. The rule is hard: no phase passes its gate without a fully successful testset run.

9.3 Source verification per entry

Every typology entry has a source-verification section. Every cited verse is checked against the strong/[book].jsonl files. The citation must match the source text byte-for-byte.

9.4 Cross-reference audit

The cross-reference audit produces three artefacts: a list of broken links, a list of asymmetric links, and a top-15 list of most-referenced entries.

9.5 Naming-discipline scan

Before publication, a final-pass scan removes all internal personal names, doctrinal-school names, and project-internal identifiers from the public artifacts.

9.6 Reproducibility verification

The reproducibility verification reruns each entry's zoekquery_gebruikt and confirms that the command yields the same verse list it yielded when the entry was written.

Appendix A. Glossary

aiōn. Greek noun meaning "age" or "eon". Concordant: "eon". Tradition often: "forever" or "world".

aiōnios. Greek adjective derived from *aiōn*. Concordant: "eonian".

attentive-reader presumption. The interpretive principle that the Bible was written for readers who knew the text well, so deliberate echoes are read as deliberate.

Block A through F. The named output blocks. A: direct reading. B: source citations. C: typology cross-references. D: concordant rendering. E: coherence check. F: short transcript.

brith chadashah. Hebrew "new covenant". Jeremiah 31:31.

concordant translation. A discipline that fixes one target word per source word across all occurrences.

ekklēsia. Greek noun for assembly. Concordant: "called-out". Tradition: "church".

hamartia. Greek noun for missing-the-mark. Tradition: "sin".

Hxxxx and Gxxxx. Strong-code prefixes. H for Hebrew or Aramaic, G for Greek.

JSONL. Line-delimited JSON, used for source-text layers.

LXX. The Septuagint, the ancient Greek translation of the Hebrew Bible.

N1, N2, N3. Certainty levels. N1: explicit. N2: inferable. N3: suggested clue.

nephesh. Hebrew noun for breath, soul, or living-being. Concordant: "soul" everywhere.

preflight. First stage of question processing: parsing, language detection, trigger matching.

seven-step protocol. The procedure that every new typology entry must pass.

sola scriptura. The operational rule that no claim may rest on anything other than the bundled source-text corpus.

Strong code. A 19th-century indexing system mapping each Hebrew, Aramaic and Greek word to a unique number.

two-witness rule. A pattern requires at least two independent witnesses in the canon.

typology. The study of patterns the source text builds across the canon.

typos. Greek noun meaning "type" or "pattern". Romans 5:14: Adam *typos tou mellontos*.

Appendix B. References and attribution

The source-text data bundled with this project is derived from open-source upstream projects.

B.1 Hebrew and Greek source text

Portions of the source text in `Kennis/puur/` derive from the publicly available interlinear data prepared by the `scripture4all` project (<https://www.scripture4all.org/>).

B.2 Strong-code annotation

The Strong-code annotated text in `Kennis/strong/` derives from `STEPBible-Data`, published by Tyndale House, Cambridge (<https://github.com/STEPBible/STEPBible-Data>).

B.3 LXX mapping

The Hebrew-to-Greek lexical mapping was constructed by aggregating co-occurrence frequencies across openly available Septuagint editions.

B.4 License obligations for redistribution

If you redistribute this project, you must preserve attribution for each upstream data source. See `ATTRIBUTION.md` in the repository root.

Appendix C. Sample workflows

Three concrete workflows.

C.1 Word-study

Question. What does *aiōn* mean across the New Testament? **Step 1.** `python3 skill/typologie_zoek.py --strong G165`. **Step 2.** 122 occurrences. **Step 3.** Cluster by meaning. **Step 4.** Read concordant master. **Step 5.** Compare against tradition.

C.2 Verse-study

Question. Walk me through Romans 5:14. The skill loads the verse, the preflight activates two typology entries, builds Blocks A through E, marks N1 because the pattern is explicit.

C.3 Pattern-study

Question. Is there a pattern around the number 120? **Step 1.** Formulate hypothesis. **Step 2.** `python3 skill/typologie_zoek.py --cijfer 120`. **Steps 3-5.** Inventory, cluster, coherence test. **Steps 6-7.** Write entry and update indexes.

End of document.